
LSTM for Text Generation and Classification

Frédéric Boileau
Jimmy Leroux
Nicolas Laliberté

frederic.boileau@umontreal.ca
jimmy.leroux@umontreal.ca
nicolas.laliberte@umontreal.ca

Abstract

In the context of our project for the IFT6269-A2018 we have investigated the different available models for the generation and classification of natural text data. While we were initially interested in fully probabilistic models such as Hidden Markov Models (HMM), a quick review of the contemporary literature on the topic of pattern recognition on sequences made it clear that neural networks (NN) provided the better toolset. In the end we implemented two Long Short-Term Memory networks, for generation and classification respectively which we trained on subsequences of some corpora of prose fiction. Despite the rather coarse nature of our implementation the generators were able to produce legible and decently structured text which reflected the material used for training; in the style of the writing for example. The classifier reported excellent metrics however there are several interesting questions as to what overfitting consists of when classifying prose fiction. The most obvious is the one of proper nouns, but others such as temporal markers are more subtle. The solutions to be decided with some context of the intent of the task, so as to be the answer to a well-defined question.

1. Introduction

Hidden Markov Models used to be the go-to probabilistic tool to reason about sequential data; the Markov assumption however proves to often be unreasonably strong. Adding links to form higher order chains is not a scalable solution as the computational complexity grows exponentially in the order of the chain.

Recurrent Neural Networks (RNN) constitute a family of neural network architectures specialized to process sequential data which can forfeit Markovian assumptions while remaining tractable. RNNs can track longer range dependencies while staying tractable by leveraging the simple idea of sharing parameters across the model (Goodfellow et al., 2016). Concretely this means adding loops to the hidden units of the neural network. RNNs have been successfully used in diverse domains for generating sequences such as music and text (Graves, 2013). Despite the aforementioned features, naive RNNs suffer from the fact that the influence of some input to the hidden layer either decays or blows up exponentially in time, a phenomenon referred to in the literature as the vanishing gradient problem. The cost of abandoning probabilistic models such as the HMM in favor of neural networks is the loss of a fully probabilistic interpretation. There has recently been an increased interest into finding reasonable probabilistic interpretations to RNNs, see for example (Choe et al., 2017). On the other hand the very existence of some monolithic notion of “interpretability” has been recently questioned, see (Lipton, 2016) for a philosophically inclined take on the question.

2. Neural Networks Architecture

2.1. RNN and the Vanishing Gradient




Figure 1. Unfolded RNN network. Figure taken from Christopher Olah’s blog.

$$h_t = \sigma(Ux_t + Vh_{t-1} + b_h) \quad o_t = \text{softmax}(Wh_t + b_o)$$

Where U , V and W are weights matrix and the vectors b are bias parameters. Consider the gradient of $o_{t+\delta}$ with respect to h_t . Applying the chain rule according to the graph above we get:

$$\nabla_{h_t} o_{t+\delta} = \left(\prod_{k=t+1}^{t+\delta} V^T \text{diag}(h_k(1-h_k)) \right) \nabla_{h_{t+\delta}} o_{t+\delta}.$$

Thus, as δ grows, the gradient grows exponentially with V . If V is small or large then the gradient will either vanish or explode. A myriad of solutions exist such as regularization through weight noise, the Long Short Term Memory architecture tries to tackle this issue on a higher level than regularization.

2.2. LSTM Architecture

To go from a RNN to a LSTM we replace hidden units with components known as memory cells. Our presentation and notation follows (Lipton, 2015).

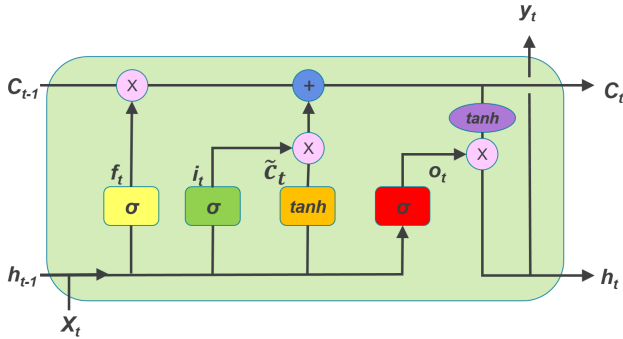


Figure 2. Memory cell of a LSTM. Figure taken from Christopher Olah’s blog.

Intuitively, RNNs have long-term memory in the form of matrix weights, they change during the training; encoding through training some general knowledge about the data. They also have short-term memory in the form of activation passing from each node to successive ones. The memory cell introduced in the LSTM model formalizes those notions and provides a framework to control their impact on the internal representation of the network.

- f_t, i_t, o_t : Respectively forget, input and output gates.
 - Sigmoidal units activated through x_t (current input) and h_{t-1} (past hidden unit output)
 - f_t controls the recursive flow

- i_t and o_t control the input and output flow respectively
- $h_t = o_t \odot \tanh(c_t)$ where \odot denotes element wise multiplication.

- $c_t = \mathbf{f}_t \odot \mathbf{c}_{t-1} + i_t \odot \tilde{c}_t$: The cell which has a self-connected edge with a fixed unit weight, thereby delegating control of recursion to the gate

The architecture of the memory cell improves the capacity to learn long time dependencies. Indeed, following the definition of the cell state, we have

$$c_t = \sum_{k=0}^t f_t \odot \dots \odot f_{k+1} \odot i_k \odot \tilde{c}_k.$$

We can see that the gradient may pass over long time gaps as long as the forget gates are open (closed to 1) which can be achieved with a appropriate initialization setting a large bias for the gates f .

3. Implementation

We now discuss training of the LSTM generators. During initial prototyping we trained on quite varied datasets to see how structure would impact the generation. The four datasets were the Harry Potter novel series, the Lord of The Rings series, a Kaggle dataset of jokes and a collection of excerpts from shakespeare. This produced reasonable synthetic data (presented below, note that the text for training is easily identifiable) and we afterwards focused exclusively on prose fiction, selecting authors which seemed close in terms of style.

The two text sequences of prose used for the subsequent training of generators were two concatenated novels of Jane Austen and of George Eliot respectively. All in English and utf-8 encoded. Punctuation and structure was left unprocessed. Each dataset was split in sequences of 50 tokens (i.e. “words”) and a dictionary was built from the complete input ($\approx 60k$ unique tokens), defining the input space for the networks. Vector encoding of this space was used through an embedding layer mapping the words to a real vector space of dimension 256. Available embeddings such as word2vec and glove were initially used but proved to be more cumbersome than our own trained version.

We trained four LSTM networks with the aforementioned data, a generator and a classifier, once without preprocessing and another where the NLTK Part of Speech Tagger (POS tagger) was used to filter out proper names and replace them with generic ones from the names dataset of the NLTK library. The generic name replacement was shared across the Jane Austen

and George Eliott datasets. This was an effort to make the classifier be trained on rather structural aspects of the prose, such as syntax, length of sentences and broad notions of style as opposed to simply directly matching vocabulary which would have been trivial with the names of the characters.

Training was achieved at “word” level (tokenized with NLTK). Character level had been previously envisioned for the rest of the project as it is more flexible and can learn new words and structural information (Graves, 2013) for the generators). Despite this we kept the training at word level for multiple reasons, the primary one being robustness towards unicode characters which might vary between versions of the text available and the second one being speed of convergence.

To generate the sequences the trained generators were initialized with a random word drawn from the dictionary and the most likely next word was fed back in the network until the desired sequence length was reached. 1000 sequences (250 per model) were generated. The classifier was trained on the original datasets and were then used to classify the generated synthetic data estimate which dataset (or model, equivalently) was used for its generation, thereby giving us some “metric” of the quality of the data.

4. Empirical Results

4.1. Models details

We used the many-to-one architecture like the one below for our classifier. We used the last output of the




Figure 3. Many-to-one architecture

LSTM as our input for the classifier, disregarding all the other outputs. The last hidden state contains information about all the sequence through the memory cell.




Figure 4. Classification error for Austen and Eliot datasets.

4.2. Results

The error we used here is the cross-entropy loss, which is standard for classification tasks. We see on our training curves (figure 2) that we have reach 96% for the classification accuracy.

For the text generation models, we used a many-to-many architecture like the one showed in the RNN section with one layer for computational reasons. LSTMs are with a cell state of 512 dimensions. We trained trough minimization of the cross-entropy loss, which is equivalent to perplexity; the standard metric for language modeling (Graves, 2013).

Examples of generated quotes:

- “ well , we ’ ll do it with a wand , ” said hermione . “ really ? ” said harry , looking at each other .
- what looked about this way , the black citadel , was far from the darkness , the ring was heard , but the sea big was big , and a great ring was in his battle .
- failure is a beginning of love and a family which comes from god .
- ” that now my mind shall screens his music , ” ” and i to give thee my vulgar heaven , ” ” i am your brother .

5. Discussion

We first tackle the results on the text generation. It seems to be hard to draw major conclusions from the measure that is Bit Per Character at this stage. The preprocessing of character names doesn’t seem to have an interpretable effect on the results of the generators




Figure 5. Classification error for Austen and Eliot with generic character names

Dataset	BPC	Perplexity
Jane Austen	1.00	33
George Eliot	1.11	47.18
Jane Austen generic names	1.29	88.03
George Eliot generic names	1.09	43.41

Dataset	BPC	Perplexity
Harry Potter	1.00	33
LOTR	1.02	35
Random quotes	1.10	45
Shakespeare	0.94	26

as it gets higher for Jane Austen but lower for George Eliot. The values for all text data tried so far were each too close to be a solid basis for interpretation.

In the case of classification, training converges very fast, with or without replacing character names with generic names. Moreover it performs as well on synthetic data, meaning when trained on the original novels it can identify which generator output which synthetic data. This leads us to think that the generator networks captured quite well the defining features of the authors. Now the important question is whether or not those features are relevant in determining authorship in general.

Authorship analysis is a field of computational linguistics where the author’s identity is to be extracted from the writing style of text. (Ding et al., 2016). It has diverse applications beyond complementing the work of traditional philologists such a fraud and plagiarism detection or even identifying cybercriminals on underground forums(Afroz et al., 2014). Classically

the field depended on statistical analysis of manually engineered features such as word frequency, length of sentences etc.

Despite its great importance and the fact that it is fundamentally quite well defined, i.e. the question who wrote this text is more well posed than the ones concerning sentiment analysis for example, there are major challenges. Clearly one needs to differentiate between writing features specific to writing style as opposed to context/ content specific features as Rosenthal and Yoon mention in their paper on the application of stylometric techniques to detect multiple authorship in legal decision (2011).

In this context the simple “anonymization” of the text sequences we performed has the intent to disentangle the contextual features from the imprint of the author. We have mentioned before that the classifier could be used as some kind of metric on the synthetic data. Conversely, suppose one is fairly confident in the quality of the generated text, i.e. after reading its content and using some domain knowledge experts determines that it reflects quite well the imprint of an author’s style. In this case the generated data could be a candidate for training of classifiers to detect authorship, when the data is sparse for instance or simply because the very nature of synthetic data could prevent overfitting to less relevant features, the author’s footprint being in some sense crystallized in the generator’s features. Of course this is purely circular if both directions are considered at the same time and we have found the latter(using synthetic data to evaluate the generalization of the classifier) to be more interesting both in the results it yields and in its theoretical formulation.

In essence the goal would be to use text generation to trim from data all contextual information that is not considered to be relevant in determining the author (such as names or temporal and spatial markers in some cases) while keeping the semantics of the text intact. Afterwards using that stripped text to train the generators to obtain more synthetic data which reflects the author in its structural semantics, beyond crude word frequency. Of course the question as to what to keep from the original text, what to trim and generally what features constitute an author’s footprint is one that for the moment relies heavily on domain-knowledge, however, given that knowledge, LSTM can leverage it to perform much deeper analyses than classical statistics provide, relying on structure and not just isolated features.

6. Conclusion and Further Work

Through our project we have explored the different available methodologies for the generation and classification of textual data. We explored both problems as they seem complementary in some respects however for practical applications classification is the most interesting and relevant.

The most capable tools right now for sequence classification and generation revolve around NN and more precisely the LSTM architecture seems to be the go to building block for those applications. One criticism of NN and deep learning based pattern recognition techniques has revolved around notions of interpretability, however as Lipton discussed in his paper (2017) there is no single monolithic notion of interpretability. One of the reasonable challenges he mentions is the one of contestability of the results of a model when we need ethical decision-making. Classical stylistic techniques rely on statistics which are fairly easy to explain and/or refute but the result from a LSTM might be harder to verify. This issue can be addressed in many ways but first the notion that results need be understandable by a reasonably educated person might seem over reaching as we already rely on hard to contest forensic techniques such as DNA testing. To address this issue instead of simply dismissing it one might consider ways to make the internal representation of a RNN model more cogent, ironically one possible avenue Lipton suggests is to train a RNN to map the internal state of the model used to a textual explanation.

One critical application of text classification is authorship analysis of which we have surveyed examples in diverse settings such as identifying cyber criminals in underground forums or determining multiple authorship in legal cases.

We have obtained surprisingly good results in both classification and text generation despite our rather coarse training. The generative models clearly learned some structural information, often matching speech marks in a reasonable way, and yielding more complex sentence structures using hypotaxis when trained from authors such as Jane Austen while using simpler shorter sentences with conjunctives in the case of popular authors such as JK Rowling.

In the case of classification for authorship analysis, overfitting seems to be difficult to prevent without a clear methodology of what constitutes contextual information and what is more intrinsic to an author's style. As discussed before there are many ways to answer this question and domain knowledge should not

be dismissed. Replacing character names in prose fiction by generic ones seems to have reduced the confidence in our model. Explorations into what other irrelevant features should be trimmed from the set is thus warranted such as spatial markers. That is, obviously, when those markers are deemed irrelevant to the task. In some cases, with larger datasets with many character names that end up being shared across corpora, their choice by an author could very well be considered to be relevant as markers of their style.

7. Citations and References

References

- Afroz, S., Islam, A. C., Stolerman, A., Greenstadt, R., and McCoy, D. Doppelgänger finder: Taking stylometry to the underground. In 2014 IEEE Symposium on Security and Privacy (SP), volume 00, pp. 212–226, May 2014. doi: 10.1109/SP.2014.21. URL doi.ieeecomputersociety.org/10.1109/SP.2014.21.
- Choe, Yo Joong, Shin, J. C., and Spencer, Neil. Probabilistic interpretations of recurrent neural networks, 2017.
- Ding, Steven H. H., Fung, Benjamin C. M., Iqbal, Farkhund, and Cheung, William K. Learning stylistic representations for authorship analysis. CoRR, abs/1606.01219, 2016. URL <http://arxiv.org/abs/1606.01219>.
- Goodfellow, Ian, Bengio, Yoshua, and Courville, Aaron. Deep Learning. The MIT Press, 2016. ISBN 0262035618, 9780262035613.
- Graves, Alex. Generating sequences with recurrent neural networks. CoRR, abs/1308.0850, 2013. URL <http://arxiv.org/abs/1308.0850>.
- Lipton, Zachary Chase. A critical review of recurrent neural networks for sequence learning. CoRR, abs/1506.00019, 2015. URL <http://arxiv.org/abs/1506.00019>.
- Lipton, Zachary Chase. The mythos of model interpretability. CoRR, abs/1606.03490, 2016. URL <http://arxiv.org/abs/1606.03490>.
- Rosenthal, Jeffrey S and Yoon, Albert H. Detecting multiple authorship of united states supreme court legal decisions using function words. The Annals of Applied Statistics, pp. 283–308, 2011.